

Accelerate Your Code

with the Accelerate framework

Jeff Biggus – HyperJeff, Inc – @hyperjeff

Who am I?

- Programmer from a science background, not from CS
- Creator of various OS X resources over the years
- Someone who thought the biggest announcement about iOS 4 was the inclusion of Accelerate.framework
- CocoaLit.com
- <http://blog.hyperjeff.net/blasLookup.pdf>

Accelerate

Accelerate Framework

vecLib

vImage

BLAS

LAPACK

vDSP

vMisc

(vForce)

Jaguar
iOS 4

Tiger
iOS 5

History



The Idea: HPC



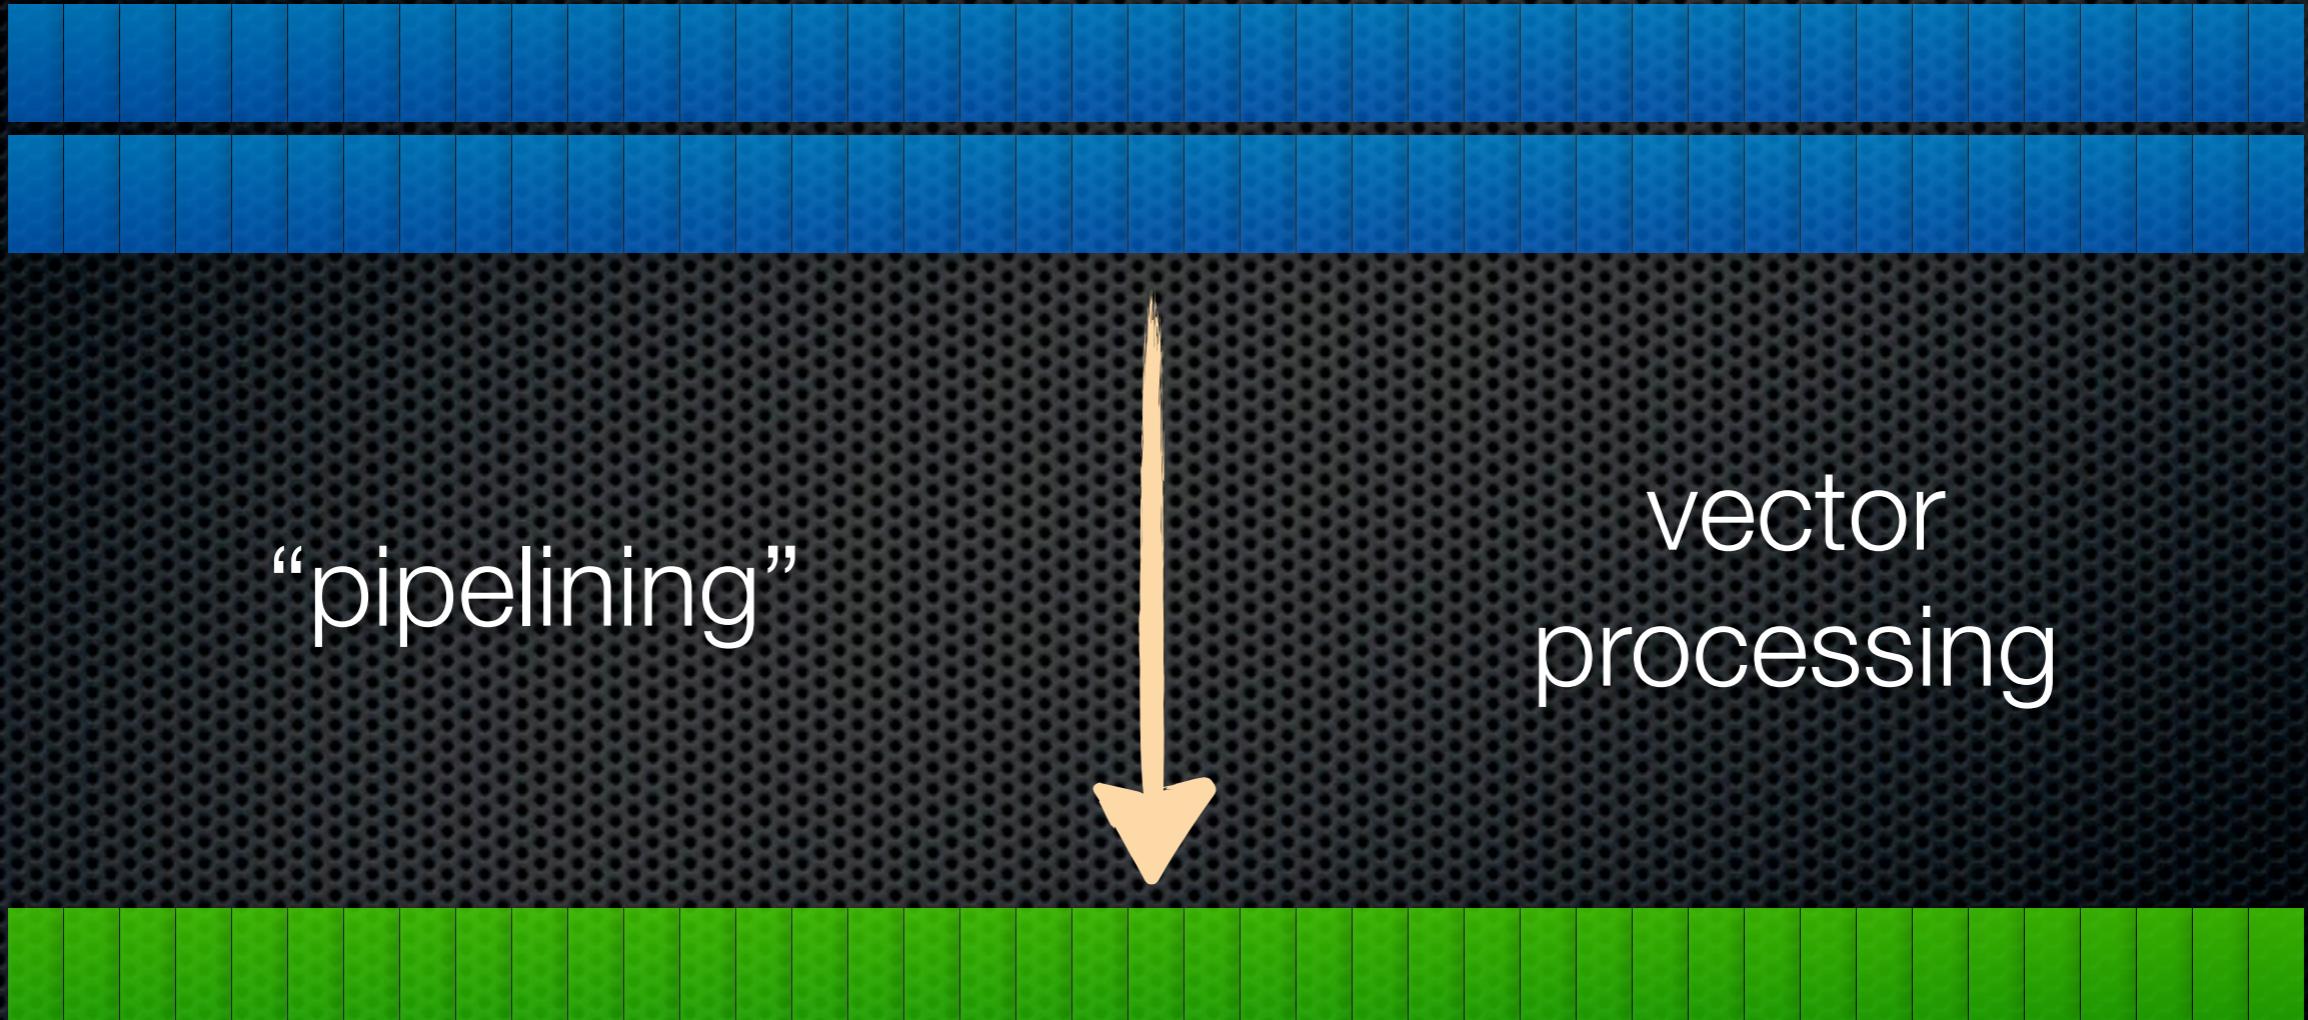
```
for (...) {
```

```
    x y → z
```

```
}
```

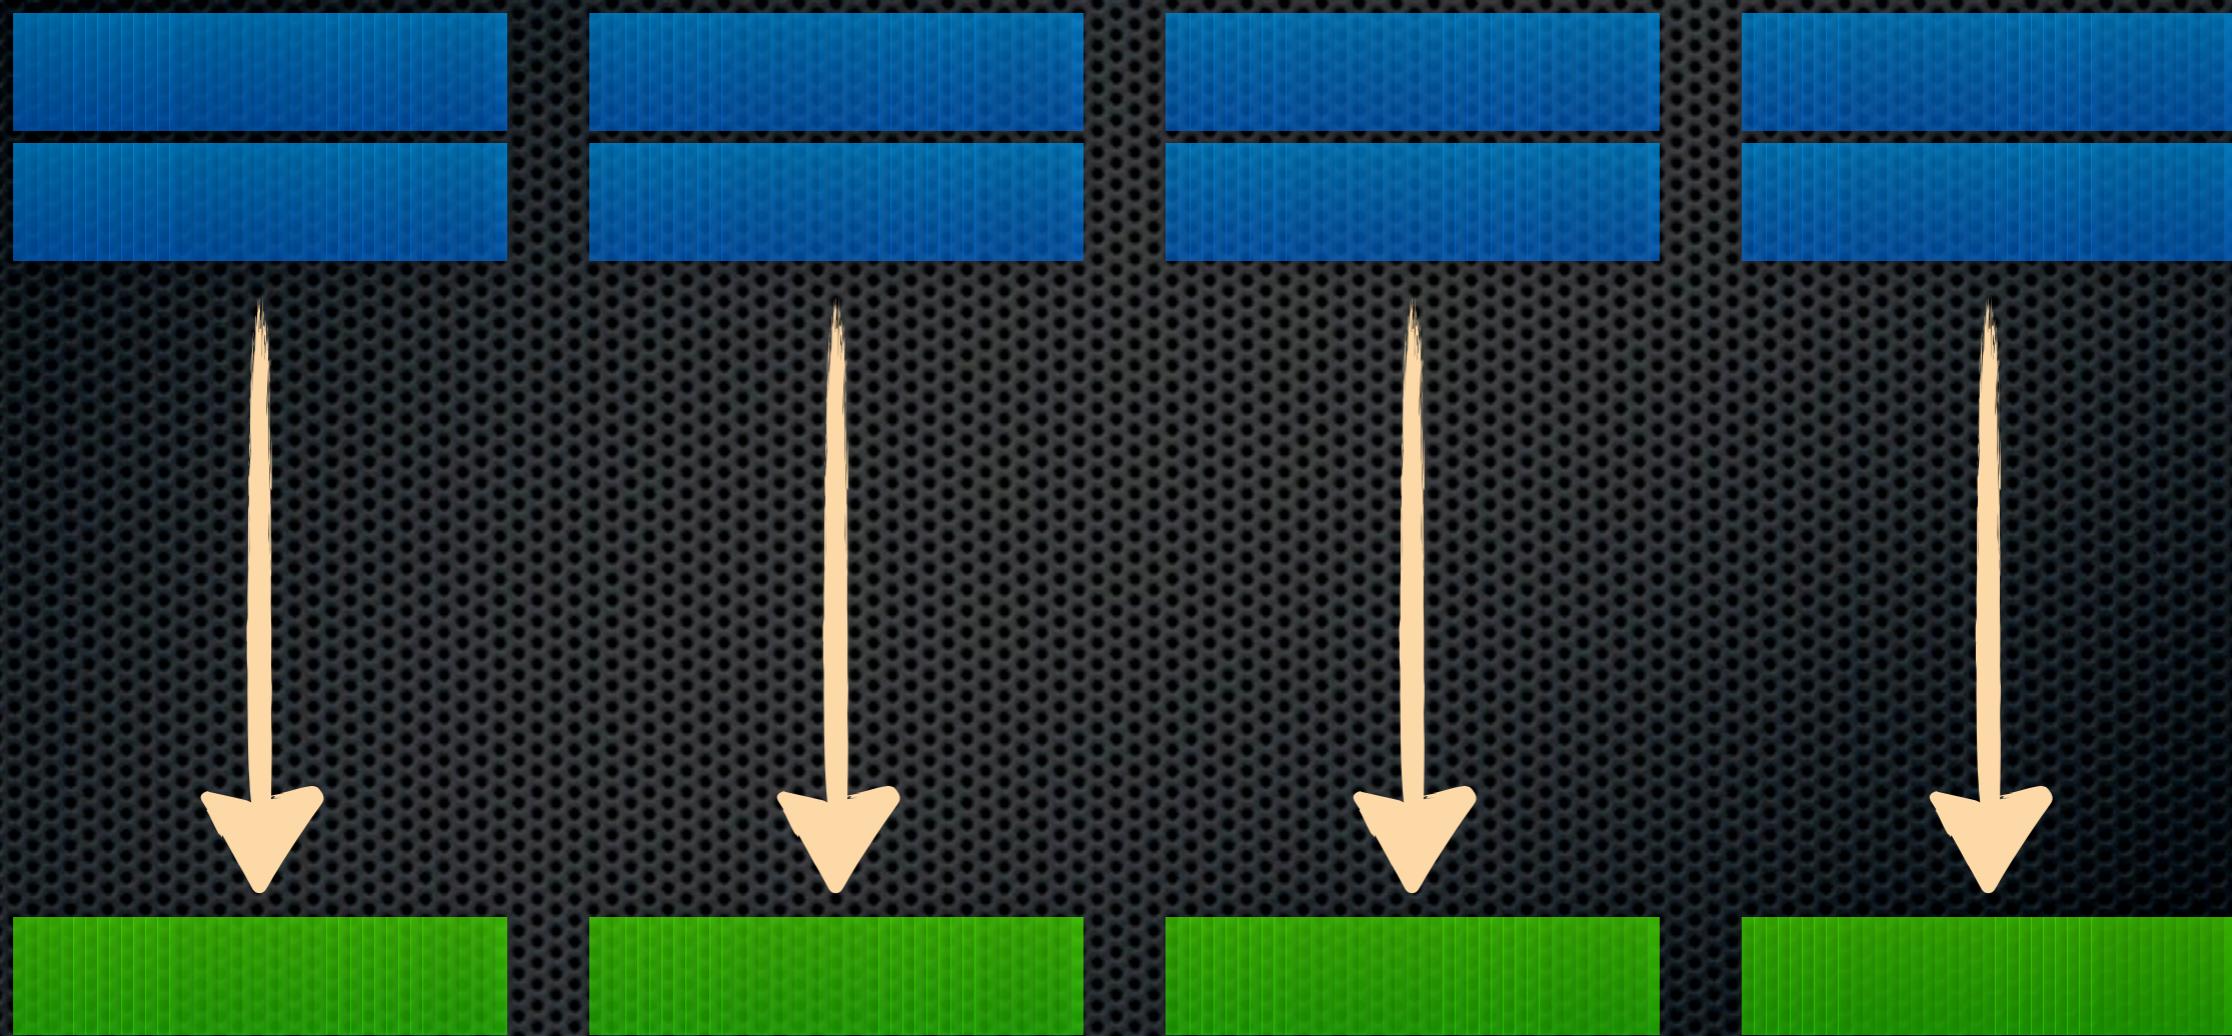


The Idea: HPC



SIMD: single-instruction, multiple-data

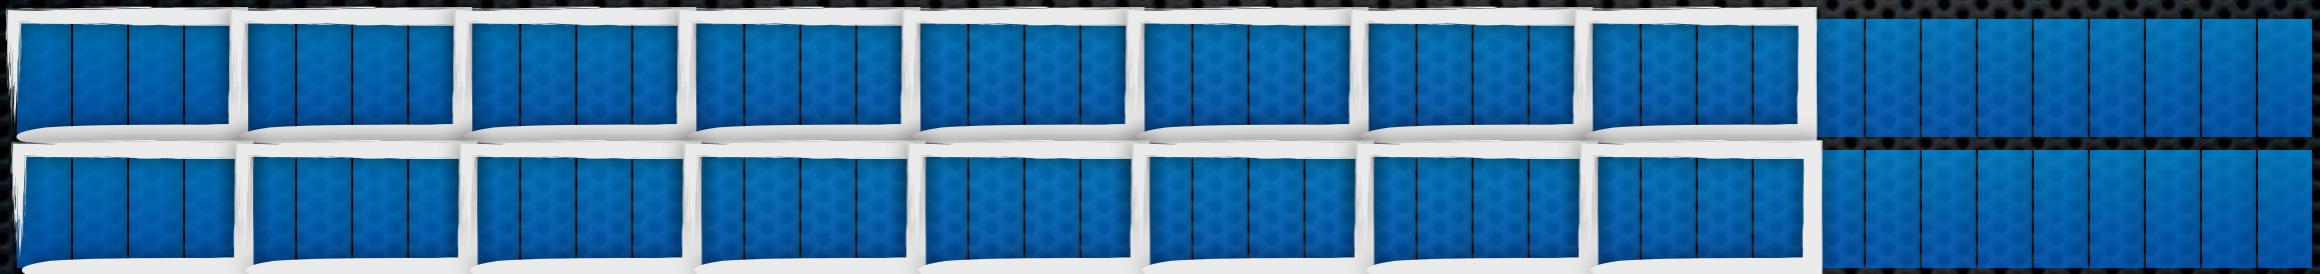
The Idea: HPC



Parallelism

■ = 1 byte

1 float
register



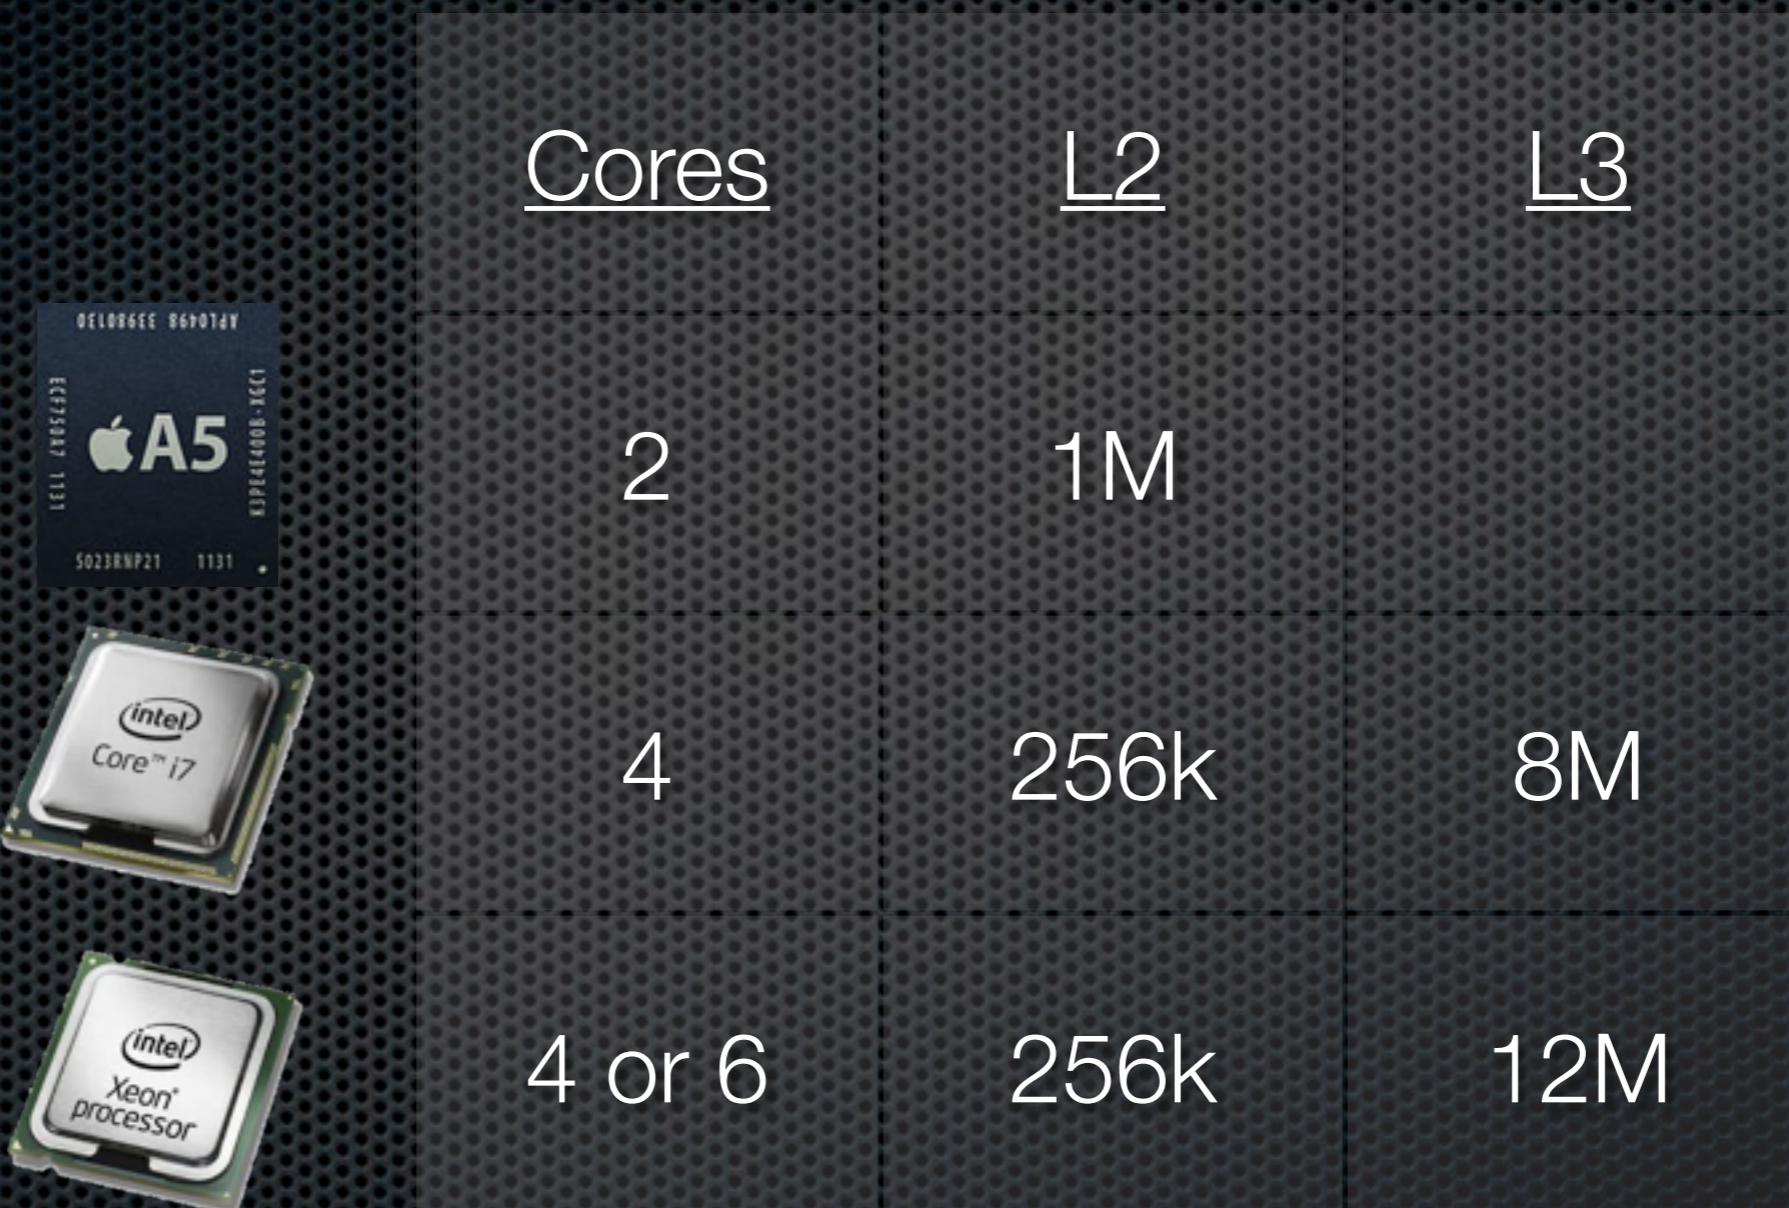
*but, we also have **16** 128-bit registers*



More efficient use of registers, more data processing per clock tic, less energy per work unit

Hardware Optimization

<u>Cores</u>	<u>L2</u>	<u>L3</u>
2	1M	
4	256k	8M
4 or 6	256k	12M



The image displays three processor components on a dark, textured background. On the far left is the Apple A5 chip, a square blue die with 'A5' printed on it. To its right is the Intel Core i7 processor, a larger silver package with 'Core i7' visible. Below the Core i7 is the Intel Xeon processor, another silver package with 'Xeon processor' written on it.

Good, Bad & the Ugly

Cons:

- Pure C
- Some Setup Work
- Documentation
- Separate Libraries
- Resources*

Pros:

- Hardware Optimized
- Speed
- Energy Efficient
- Clean, Readable Code
- Code Re-Use

Competitive Advantage

Used in (only) ~100 GitHub projects

That's out of >500k ObjC projects!

Your competitive advantage

sometimes as simple as:

```
#import <Accelerate/Accelerate.h>
```

or

```
clang ... -framework Accelerate
```

Accelerate

- vImage 218 functions
 - vDSP 401 functions
 - vForce 80 functions
 - LAPACK 1471 functions
 - BLAS 196 functions
- 2366 functions

BLAS

BLAS

LAPACK

vForce

vDSP

vImage

- Basic Linear Algebra Subprograms
- Vector → Vector operations BLAS 1, $O(n)$
- Matrix / Vector → Vector BLAS 2, $O(n^2)$
- Matrix / Vector → Matrix BLAS 3, $O(n^3)$
- Dense matrix routines (along with LAPACK)

BLAS

LAPACK

vForce

vDSP

vImage

```
#include <Accelerate/Accelerate.h>
#include <stdio.h>

int main(int argc, const char *argv[]) {
    float x[] = { 1., 2., 3. };
    float y[] = { 3., 4., 5. };

    //      y = 10 x + y
    cblas_saxpy( 3, 10., x, 1, y, 1 );

    printf( "\n y = { %2.f, %2.f, %2.f}\n", y[0], y[1], y[2] );

    return 0;
}
```

```
% clang -o cblas_saxpy cblas_saxpy.c -framework Accelerate
% ./cblas_saxpy
```

```
y = { 13, 24, 35 }
```

BLAS

```
float x[] = { 1., 2., 3. };  
float y[] = { 3., 4., 5. };
```

```
cblas_saxpy( 3, 10., x, 1, y, 1 );
```

y: { 13, 24, 35 }

LAPACK

vForce

vDSP

vImage

```
float x[] = { 1., 9., 2., 9., 3., 9. };  
float y[] = { 3., 9., 4., 9., 5., 9. };
```

```
cblas_saxpy( 3, 10., x, 2, y, 2 );
```

y: { 13, 9, 24, 9, 35, 9 }

BLAS

LAPACK

vForce

vDSP

vImage

cblas_saxpy

cblas_

s

axpy

BLAS

type

function

[BLAS](#)[LAPACK](#)[vForce](#)[vDSP](#)[vImage](#)

```
double complex u[] = { -3. + 4.*I,  5. + 7.*I };  
double complex w[] = {  1. + 2.*I, -1. + 5.*I };  
double complex alpha[] = { 10. + 100.*I };  
  
//           w = alpha   u   + w  
cblas_zaxpy( 2, alpha, u, 1, w, 1 );
```

```
float a[] = {
    10., 5., 3.,
    5., 4., 2.,
    3., 2., 1.
};
```

```
float b[] = {
    1., 2.,
    3., 4.,
    5., 6.
};
```

```
float c[9];
```

```
cbblas_ssymm(
    CblasRowMajor,
    CblasLeft,
    CblasUpper,
    3, 2,
    1., a, 3, b, 2,
    0., c, 2
);
```

cbblas_ssymm

Multiplies a matrix by a symmetric matrix (single-precision).

```
void cbblas_ssymm (
    const enum CBLAS_ORDER Order,
    const enum CBLAS_SIDE Side,
    const enum CBLAS_UPLO Uplo,
    const int M,
    const int N,
    const float alpha,
    const float *A,
    const int lda,
    const float *B,
    const int ldb,
    const float beta,
    float *C,
    const int ldc
);
```

$$C = A \cdot B$$

C:

	40	58
	27	38
	14	20

```
float x[] = { 1., 2., 3. };

float A[] = { // upper-triangular _packed_ matrix
    2., 5., 10.,
    1., 5.,
    3.
};

// x = A x

cblas_stpmv(
    CblasRowMajor, CblasUpper, CblasNoTrans, CblasNonUnit,
    3, A, x, 1
);

% ./cblas_stpmv
x = { 42, 17, 9 }
```

actual matrix:

3	2	1	.	.	.
2	3	2	1	.	.
1	2	3	2	1	.
.	1	2	3	2	1
.	.	1	2	3	2
.	.	.	1	2	3

```
float bandedMatrix2[] = {  
    0., 0., 3., 2., 1.,  
    0., 2., 3., 2., 1.,  
    1., 2., 3., 2., 1.,  
    1., 2., 3., 2., 1.,  
    1., 2., 3., 2., 0.,  
    1., 2., 3., 0., 0.,  
};  
  
float symmetric[] = {  
    3., 0., 0., 0., 0., 0.,  
    2., 3., 0., 0., 0., 0.,  
    1., 2., 3., 0., 0., 0.,  
    0., 1., 2., 3., 0., 0.,  
    0., 0., 1., 2., 3., 0.,  
    0., 0., 0., 1., 2., 3.  
};  
  
// CblasLower
```

<http://blog.hyperjeff.net/blasLookup.pdf>

Vector to Scalar

s = float, d = double, c = complex float, z = complex double float		
sd —	$\diamond\text{dot}$	$\vec{y} \cdot \vec{x}$
-d —	$\diamond\text{sdot}$	$\vec{y} \cdot \vec{x}$
s—	$\diamond\text{dsdot}$	$\alpha + \vec{x} \cdot \vec{y}$
e—	$\diamond\text{g2dot}$	$\alpha + \vec{x}_e \cdot \vec{y}$

Vector to Vector

sd cz	$\diamond\text{copy}$	$\vec{y} \leftarrow \vec{x}$	could be used to go from strided to/from unstrided
sd cz	$\diamond\text{swap}$	$\vec{y} \leftrightarrow \vec{x}$	differences in strides cause crazy results beyond stride _{min}
sd cz	$\diamond\text{set}$	$\vec{x} \leftarrow \alpha_0 \vec{I}$	
sd cz	$\diamond\text{scal}$	$\vec{x} \leftarrow \alpha \vec{x}$	alpha complex
— (cs)(zd)	$\diamond\text{scal}$	$\vec{x} \leftarrow \alpha \vec{x}$	alpha real
— (ca)(sa)	$\diamond\text{scal}$	$\vec{x} \leftarrow \alpha \vec{x}$	cs = complex float, zd = complex double float
— (ca)(sa)	$\diamond\text{scal}$	$\vec{x} \leftarrow \alpha \vec{x}$	ca = complex float sa = complex double float

BLAS

LAPACK

vForce

vDSP

vImage

LAPACK

- Linear Algebra PACKage
 - Primarily for *solving* systems of equations
 - Compiled from Fortran (CLAPACK), so slightly odder function call issues
 - No docs for you! (in Xcode. see clapack.h – psyche!)
 - Use the netlib.org docs (some nice things)
 - www.netlib.org/lapack/explore-html

LAPACK

BLAS

LAPACK

vForce

vDSP

vImage

- Implications of Fortran underbelly
 - Matrices are all column-major-ordered
 - All functions are postfix with _
 - *All* arguments sent to functions are references only

BLAS

LAPACK

vForce

vDSP

vImage

$$A \times = b$$

```
float A[] = {  
    3., 1., 3.,  
    1., 5., 9.,  
    2., 6., 5.  
};  
  
float b[] = { -1., 3., -3. };  
  
int output, pivot[3], numberofEquations = 3,  
    bSolutionCount = 1,  
    leadingDimA = 3, leadingDimB = 3;  
  
sgesv_( &numberofEquations, &bSolutionCount,  
    A, &leadingDimA, pivot, b, &leadingDimB, &output );
```

```
float A[] = {  
    3., 1., 3.,  
    1., 5., 9.,  
    2., 6., 5.  
};  
  
float b[] = { -1., 3., -3. };  
  
int output, pivot[3];  
  
sgesv_( &(int) {3}, // number of equations  
        &(int) {1}, // solution count  
        A,  
        &(int) {3}, // leading dimension of A  
        pivot, b,  
        &(int) {3}, // leading dimension of b  
        &output );
```

vForce

vDSP

vImage

BLAS

LAPACK

vForce

vDSP

vImage

```
n = 6
A = zeros((n,n))
x = zeros(n)

for i in range(n):
    x[i] = i/2.0
    for j in range(n):
        A[i,j] = 2.0 + float(i+1)/float(j+i+1)

b = dot( A, x )
y = linalg.solve( A, b )
```



```
int n = 6, output, pivot[n], bSolutionCount = 1;
double A[n][n], x[n], b[n];
```

```
for (int i=0; i<n; i+=1.) {
    x[i] = (float)i / 2.;
    for (int j=0; j<n; j++)
        A[j][i] = 2.0 + (float)(i+1)/(float)(j+i+1);
}
```



```
cblas_dgemv( CblasColMajor,CblasNoTrans, n,n,1,A,n,x,1,0,b,1 );
dgesv_( &n, &bSolutionCount, A, &n, pivot, b, &n, &output );
```

```

n = 6
A = zeros((n,n))
x = zeros(n)

for i in range(n):
    x[i] = i/2.0
    for j in range(n):
        A[i,j] = 2.0 + float(i+1)/float(j+i+1)

b = dot( A, x )
y = linalg.solve( A, b )

```

BLAS

LAPACK

vForce

vDSP

vImage

```

int n = 6;
Matrix *A = [ [Matrix alloc] initWithDimensions:(mIndex){ n, n }];
Vector *x = [ [Vector alloc] initWithDimension:n];

for (int i=0; i<n; i++) {
    [x setValue:(float)i / 2. atIndex:i];
    for (int j=0; j<n; j++)
        [A setValue:2.0 + (float)(i+1)/(float)(j+i+1)
            atIndex:(mIndex){ i, j }];
}

Vector *b = [A vectorMultiply:x];
Vector *y = [A linearSolve:b];

```

o.Ô



BLAS

LAPACK

vForce

vDSP

vImage

```
@interface Vector : NSObject  
@property (assign) double *v;  
@property (assign) int dimension;  
//...
```

```
@interface Matrix : NSObject  
@property (assign) double *A;  
@property (assign) int m, n;  
//...
```

```
- (Vector *)vectorMultiply:(Vector *)x {  
    Vector *y = [[Vector alloc] initWithDimension:n];  
    cblas_dgemv( CblasColMajor, CblasNoTrans,  
                m, n, 1, A, m, x.v, 1, 0, y.v, 1 );  
    return y;  
}  
  
- (Vector *)linearSolve:(Vector *)x {  
    int bSolutionCount = 1, pivot[n], output;  
    Vector *y = [[Vector alloc] initWithDimension:n];  
    cblas_dcopy( n, x.v, 1, y.v, 1 );  
    dgesv_( &n, &bSolutionCount, A, &n, pivot, y.v, &n, &output );  
    return y;  
}
```

BLAS

LAPACK

vForce

vDSP

vlImage

vForce

- Just pure mathy goodness
- v v (function name) f
- ~30 functions at the core (vs 80)

acos

cos

cosisin

cospi

asin

sin

sincos

sinpi

atan(2)

tan

tanpi

acosh

cosh

asinh

sinh

atanh

tanh

ceil

copysign

div

exp(2,m1)

fabs

floor

fmod

int

log(10,1p,2,b)

extafter

pow

rec

remainedr

(r)sqrt

[BLAS](#)[LAPACK](#)[vForce](#)[vDSP](#)[vImage](#)

vForce

vvcosf

For each single-precision array element, sets *y* to the cosine of *x*.

```
void vvcosf (
    float *,
    const float *,
    const int *
);
```

Availability

Available in iOS 5.0 and later.

Declared In

vForce.h

```
/* Set y[i] to the cosine of x[i], for i=0,...,n-1 */
void vvcosf (float * /* y */, const float * /* x */, const int * /* n */)
void vvcos (double * /* y */, const double * /* x */, const int * /* n */)
```

BLAS

LAPACK

vForce

vDSP

vlImage

vForce

y, x, n ↯ ?

$$y_i \leftarrow f(x_i) \mid 0 \leq i < n$$

z, y, x, n

$$z_i \leftarrow f(y_i, x_i) \mid 0 \leq i < n$$

sometimes C, x, n

vDSP

BLAS

LAPACK

vForce

vDSP

vImage

- Digital Signal Processing
- Audio, image, but really whatever
- Fast Fourier Transforms (1-d, 2-d)
- Vector → Scalar / Vector
- Has its own struct for complex numbers (2 kinds, even)

BLAS

LAPACK

vForce

vDSP

vImage

	1/4	
1/4	X	1/4
	1/4	

 $X = \text{ave(sides)}$

```
for (int i=0; i<GRID_SIZE; i++)
    for (int j=0; j<GRID_SIZE; j++)
        result[i*GRID_SIZE+j] = 0.25 * (
            grid[(i+1)*GRID_SIZE + j] +
            grid[(i-1)*GRID_SIZE + j] +
            grid[i*GRID_SIZE      + j-1] +
            grid[i*GRID_SIZE      + j+1]
        );
    
```

BLAS

LAPACK

vForce

vDSP

vImage

	1/4	
1/4	X	1/4
	1/4	

$$X = \text{ave(sides)}$$

```
float filter[] = {  
    0., 0.25, 0.,  
    0.25, 0., 0.25,  
    0., 0.25, 0.  
};
```

```
vDSP_f3x3( grid, GRID_SIZE, GRID_SIZE, filter, result );
```

BLAS

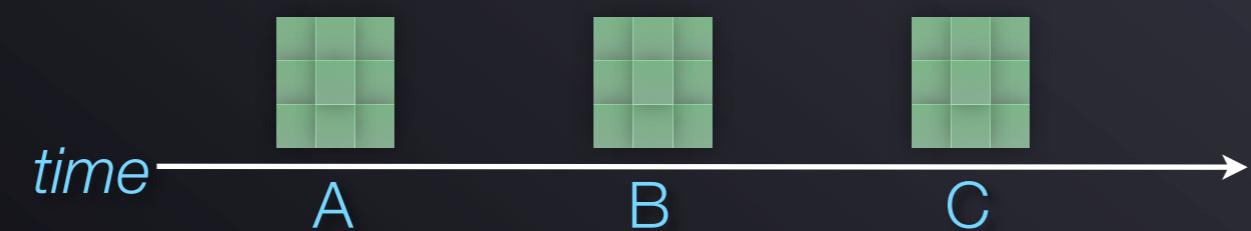
LAPACK

vForce

vDSP

vImage

	α	
α	$2-4\alpha$	α
	α	



- $\text{grid}_A + \text{conv}(\text{grid}_B) \rightarrow \text{grid}_C$

```
for (int i=0; i<GRID_SIZE; i++)
    for (int j=0; j<GRID_SIZE; j++)
        result[i*GRID_SIZE+j] =
            (2.- 4.*a)*grid[i*GRID_SIZE+j] +
            a*grid[(i+1)*GRID_SIZE+j] + a*grid[(i-1)*GRID_SIZE+j] +
            a*grid[i*GRID_SIZE+j-1] + a*grid[i*GRID_SIZE+j+1];

for (int i=0; i<GRID_SIZE; i++)
    for (int j=0; j<GRID_SIZE; j++)
        result[i*GRID_SIZE + j] -= oldGrid[i*GRID_SIZE + j];

memcpy( oldGrid, grid, GRID_AREA * sizeof( float ) );
memcpy( grid, result, GRID_AREA * sizeof( float ) );

for (int i=0; i<GRID_SIZE; i++)
    for (int j=0; j<GRID_SIZE; j++)
        points[3 * (i * GRID_SIZE + j) + 1] = grid[i*GRID_SIZE + j];
```

α	α	α
α	$2-4\alpha$	α
	α	

```
float filter[] = {
    0.0,           alpha,           0.0,
    alpha,         2.0 - 4.0*alpha, alpha,
    0.0,           alpha,           0.0
};
```

```
vDSP_f3x3( grid, GRID_SIZE, GRID_SIZE, filter, result );
cblas_saxpy( GRID_AREA, -1.0, oldGrid, 1, result, 1 );
```

```
cblas_scopy( GRID_AREA,      grid, 1, oldGrid, 1 );
cblas_scopy( GRID_AREA, result, 1,      grid, 1 );
```

```
cblas_scopy( GRID_AREA, grid, 1, &points[1], 3 );
```

BLAS

LAPACK

vForce

vDSP

vImage

```
for (int bufCount=0; bufCount<ioData->mNumberBuffers; bufCount++) {  
    AudioBuffer buf = ioData->mBuffers[bufCount];  
    int currentFrame = 0;  
    while ( currentFrame < inNumberFrames ) {  
        for (int currentChannel=0;  
             currentChannel<buf.mNumberChannels;  
             currentChannel++) {  
  
            memcpy( &sample, buf.mData+(currentFrame*4)+(currentChannel*2),  
                    sizeof(AudioSampleType) );  
  
            sample *= sin( effectState->sinePhase * M_PI * 2 );  
  
            memcpy( buf.mData+(currentFrame*4)+(currentChannel*2), &sample,  
                    sizeof(AudioSampleType) );  
  
            effectState->sinePhase += 1.0 / (asbd.mSampleRate / sineFrequency);  
  
            if (effectState->sinePhase > 1.0)  
                effectState->sinePhase -= 1.0;  
        }  
        currentFrame++;  
    }  
}
```

BLAS

LAPACK

vForce

vDSP

vImage

```
if (!setup) {  
    blockSize = inNumberFrames * buf.mNumberChannels * 8;  
  
    int blockCount[1];  
    blockCount[0] = blockSize;  
    float rampStart[] = { 0. };  
    float rampInc[1];  
    rampInc[0] = 2. / (asbd.mSampleRate / sineFrequency);  
  
    vDSP_vramp( rampStart, rampInc, bigsine, 1, blockSize );  
    vvsinpif( bigsine, bigsine, blockCount );  
  
    setup = YES;  
}  
  
vDSP_vflt16( bufferData, 1, input, 1, blockSize );  
  
vDSP_vmul( input, 1, bigsine, 1, output, 1, blockSize );  
  
vDSP_vfixr16( output, 1, bufOut, 1, blockSize );  
memcpy( bufferData, bufOut, blockSize );
```

BLAS

LAPACK

vForce

vDSP

vImage

```
switch (effectNumber) {  
  
case 0: vDSP_vclip(input, 1, clipLow, clipHigh, output, 1, blockSize);  
case 1: vDSP_vmul (input, 1, bigsine, 1, output, 1, blockSize);  
case 2: vDSP_vabs (input, 1, output, 1, blockSize);  
  
case 4: {  
    slideAdjusted = slide / 300.;  
    vDSP_vrampmul( input, 1, start, &slideAdjusted, output, 1, blockSize );  
    break;  
}  
case 5: {  
    cblas_saxpy( blockSize, slide, previous, 1, input, 1 );  
    memcpy( previous, input, blockSize );  
    break;  
}  
case 6: {  
    memcpy( temp, input, blockSize );  
    cblas_saxpy( blockSize, slide, previous, 1, input, 1 );  
    memcpy( previous, temp, blockSize );  
    break;  
}  
}
```

BLAS

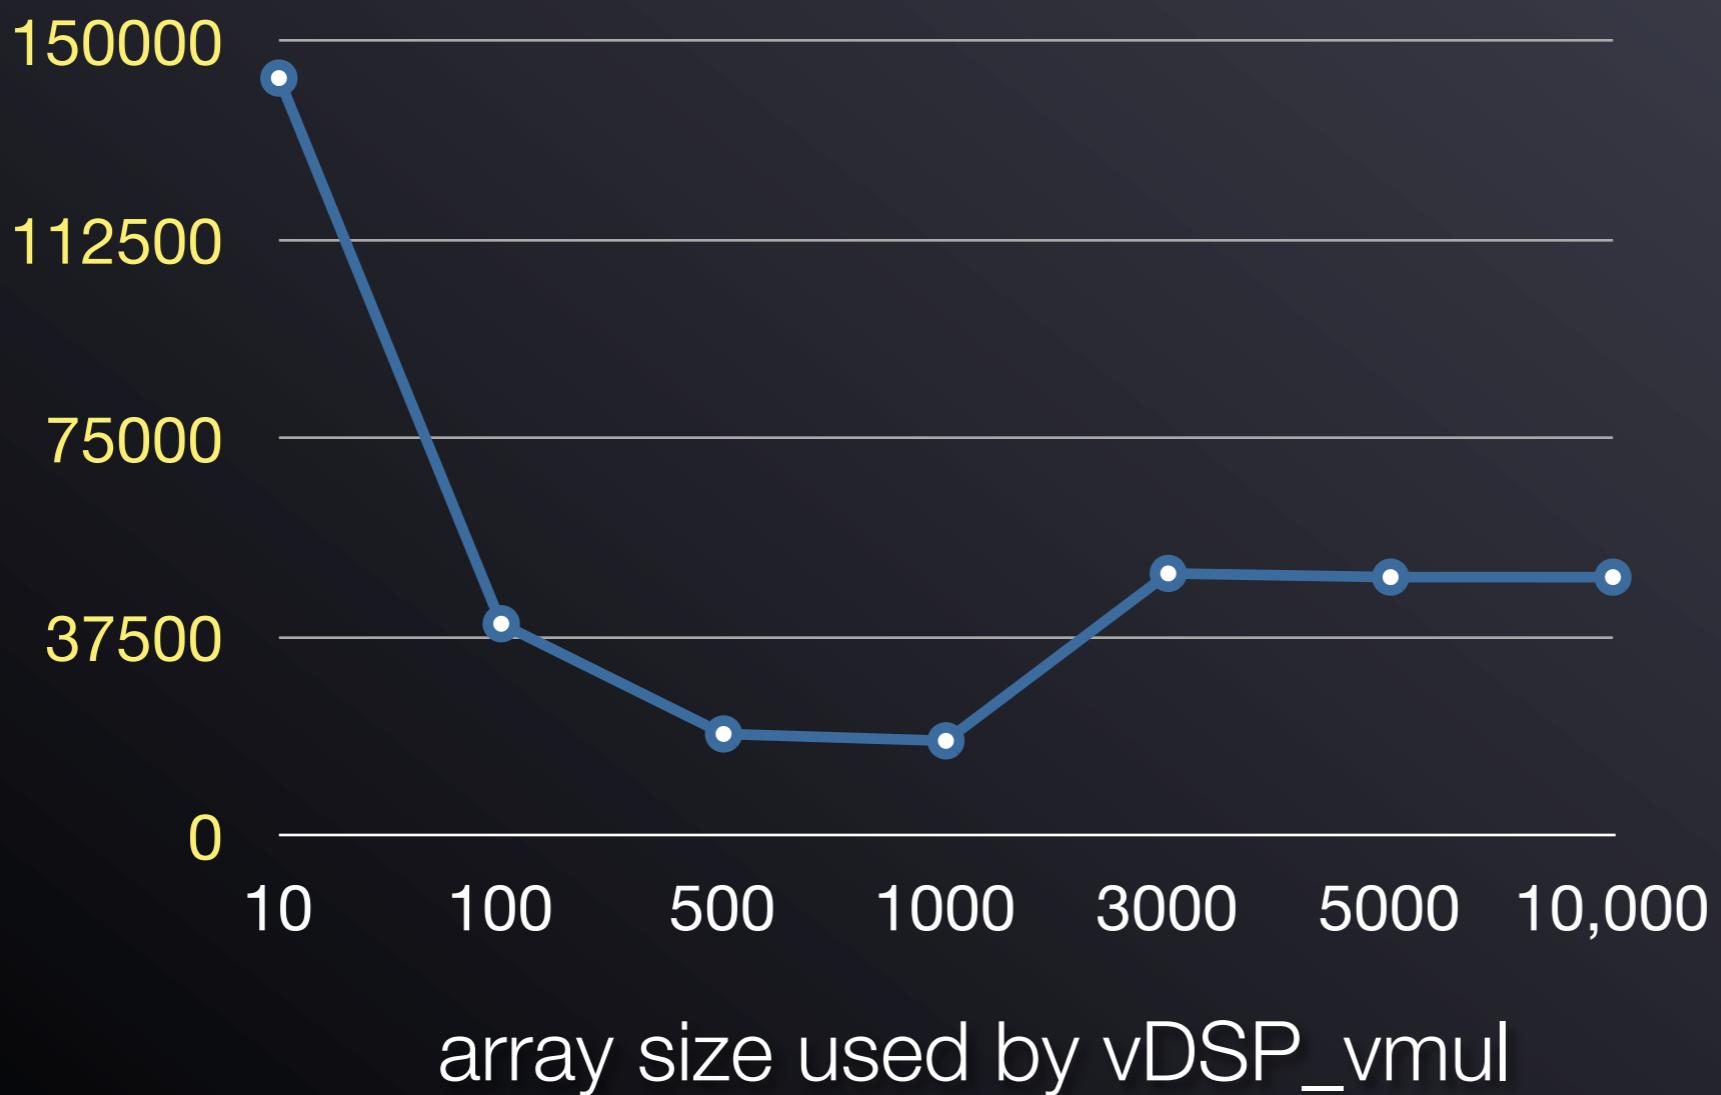
LAPACK

vForce

vDSP

vlImage

vDSP_vmul – μ s for 10,000 iterations for 10,000 element arrays



BLAS

LAPACK

vForce

vDSP

vImage

vDSP_vmul – μs for 10,000 iterations for 10,000 element arrays



BLAS

LAPACK

vForce

vDSP

vImage

vImage

- ▀ Aimed at following scenarios:



Real-time processing

- ▀ Otherwise use Core Image for regular images
- ▀ No vImage example code per se

BLAS

LAPACK

vForce

vDSP

vlImage

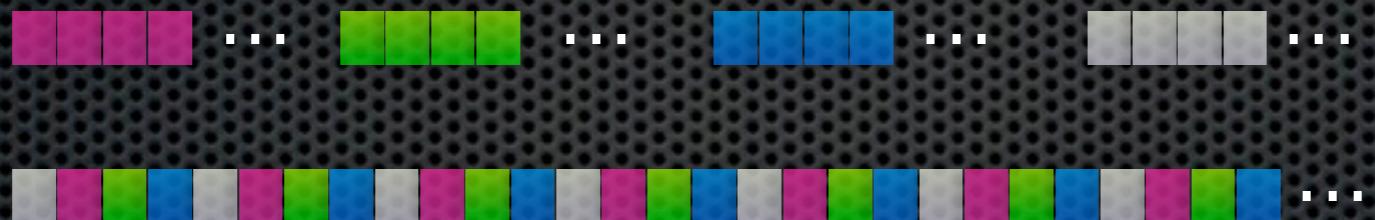
vlImage

- Shopping for image effects
- vlImage function_ format
 - ex: vlImageVerticalReflect_ARGBFFFF(...)
- Essentially ~58 functions (vs 218)
- Lots of conversion functions between vlImage formats
- vlImage_Error

vlImage

- ❖ **format**

- ❖ monochromatic “planar” vs interleaved (ARGB)



- ❖ 8-bit vs 32 (unsigned chars vs floats)
- ❖ ..._Planar8, ..._PlanarF
- ❖ ..._ARGB8888, ..._ARGBFFFF
- ❖ planar can be significantly faster

BLAS

LAPACK

vForce

vDSP

vImage

vImage

```
typedef struct vImage_Buffer
{
    void *data;
    vImagePixelCount height;
    vImagePixelCount width;
    size_t rowBytes;
}
vImage_Buffer;
```

- “Caution: except where otherwise documented, most vImage functions do not work correctly in place”

vImage

BLAS

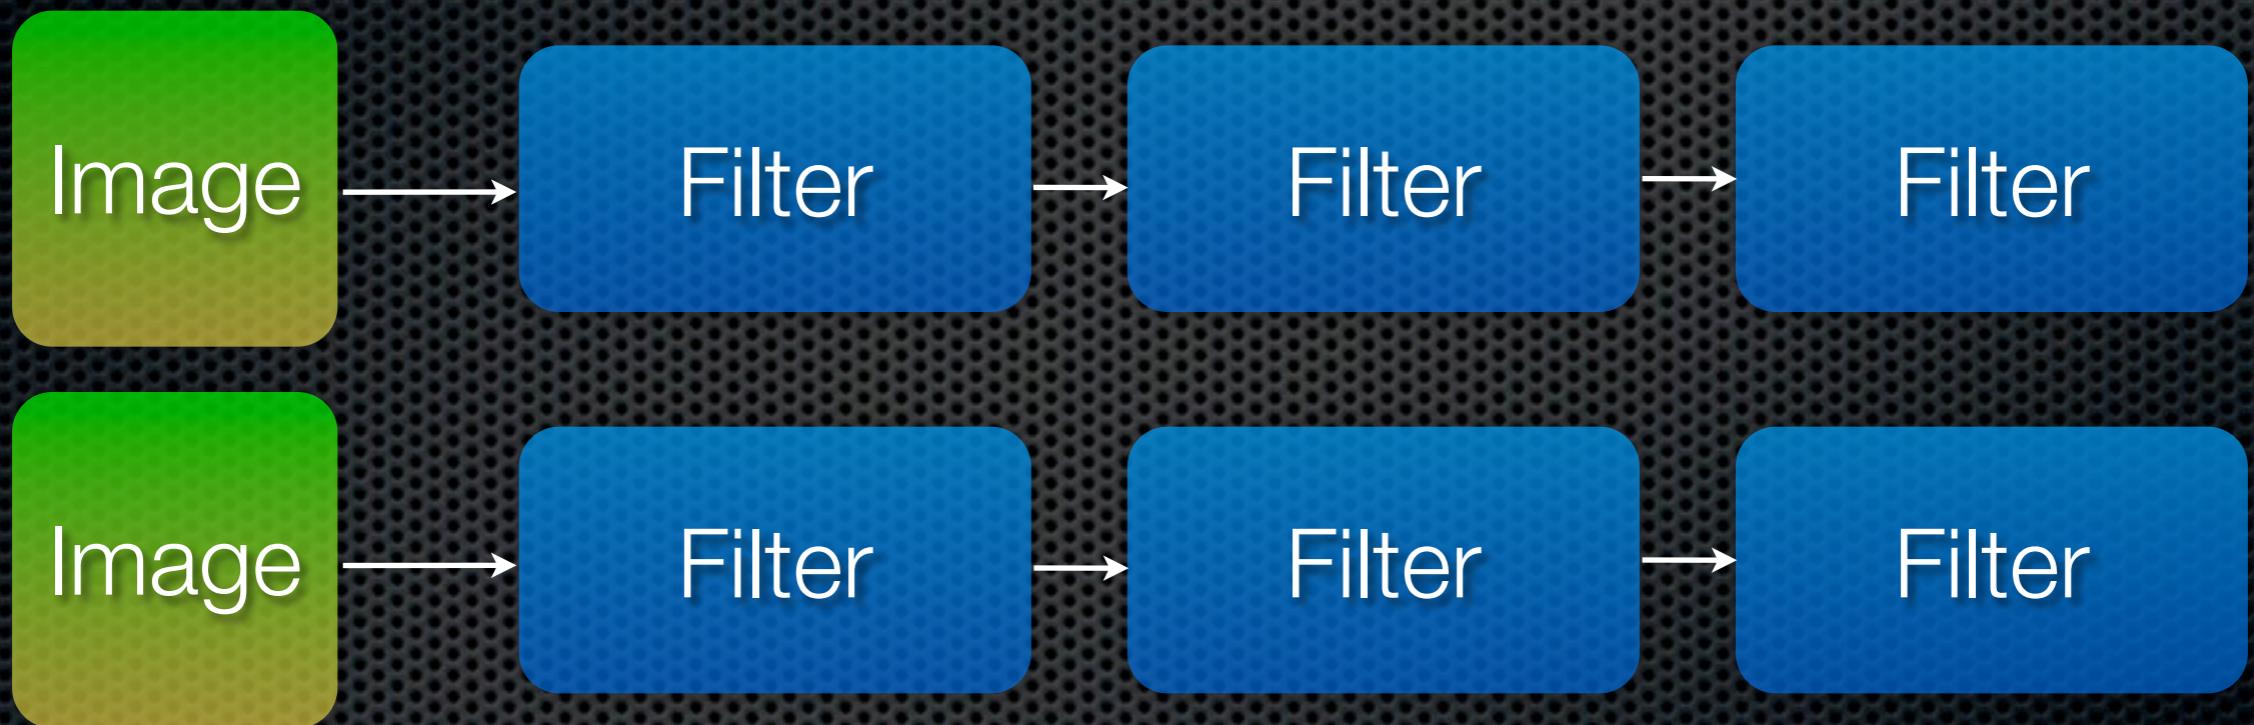
LAPACK

vForce

vDSP

vImage

L2: 256k (Xeon, i7), 1M (A5)



- “All vImage functions are thread safe and may be called reentrantly.”
- All threading done via GCD & can be turned off

BLAS

LAPACK

vForce

vDSP

vImage

vImage

- Pixel transformations
- Expansion, Contraction
- Rotation, Scaling, Warping, Reflection, Shearing
- Convolution: Smoothing, Sharpening
- Alpha compositing
- Colorspace manipulations

Accelerate Performance

- Not always the fastest solution
 - Small vectors / matrices sometimes are worse
 - Processing large data sets at once is the ideal
 - At some point, too much is also bad (faulting)
- Test out different sized data sets if it's an option
 - Suggested size of ~ L2 cache
- Avoid striding where possible

Accelerate Your Code

- Look for places where you are dealing with...
 - image effects (real-time, large)
 - audio processing
 - arrays of data of any kind needing processing
 - loops that could be turned into arrays
 - functions that are hard to do on your own
 - fns that may benefit from hardware acceleration
- Mix and match fns from any of the Accelerate libraries

Accelerate Your Code

Step back and glance at your code every once in a while and see if it couldn't make use of some bit of love from Accelerate

Check the libraries to see if there aren't some gems in there that could help make your program stand out from the crowd